

Representing Data

Dov Kruger

Department of Electrical and Computer Engineering
Rutgers University

October 17, 2024



Digital Representation of Information

- This module describes how to represent data as bits
- By the end you will be able to
 - Represent unsigned and signed Integers
 - Convert binary to/from decimal
 - Convert binary to/from hexadecimal
 - Compute twos complement numbers
 - Identify when integers overflow
 - Identify how to store text on computers
 - Identify how bits are interpreted as floating point



- With binary data, it's all context
- It is not obvious what bits represent
- In this module you will learn how the computer represents
- whole numbers, fractions, and text



- Bit: Smallest unit of data (0,1)
- Byte: Group of 8 bits
- Type of data completely depends on context



Combinations

- 2 bits: 00, 01, 10, 11
- 3 bits: 000, 001, 010, 011, 100, 101, 110, 111
- n bits = 2^n combinations
- 8 bits = $2^8 = 256$ combinations
- 16 bits = $2^{16} = 65536$ combinations
- 32 bits = $2^{32} = 4294967296$ combinations



Different Binary Data Representations

- Unsigned Integers
- Signed Integers
- ASCII text
- Unicode
- Fixed Point Fractions
- Floating Point



Integer Representations

- Integers come in different sizes (8, 16, 32, 64 bits)
- Signed/Unsigned
- Little-endian/Big-endian



Representing 3 bit Integers

bits	unsigned	signed
000	0	0
001	1	1
010	2	2
011	3	3
100	4	-4
101	5	-3
110	6	-2
111	7	-1

What happens if we add 1 to 7?

What happens if we subtract 1 from 0?



Twos Complement Arithmetic

- Consider just 8 bit number for simplicity
- First bit is sign 0 = positive, 1 = negative
- To negate a number
 - invert all bits
 - add 1
 - the resulting number is the negative of the original
 - Example: $5 = 00000101 \rightarrow 11111010 \rightarrow 11111011 = -5$
 - Example: $17 = 00010001 \rightarrow 11101110 \rightarrow 11101111 = -17$
 - Example: $-11 = 11110101 \rightarrow 00001010 \rightarrow 00001011 = 11$
 - Example:
 $-128 = 10000000 \rightarrow 01111111 \rightarrow 00000000 = -128$



Overflow and Underflow

- Overflow is when the result of a computation is too large to fit
- Underflow is the same in the negative direction
- Example: given 3-bit unsigned
 - $3 + 5 = 1000 = 8 = 000 = 0$
 - $4 + 6 = 1010 = 10 = 010 = 2$
 - $4 - 5 = 111 = 7$
- Example: given 3-bit signed
 - $3 + 2 = 101 = -2$
 - $2 - 3 = 111 = -1$
 - $3 + 1 = 100 = -4$



Overflow and Underflow

- When a result is too large, store only the low n bits
- Example: 3 bits
 - $3 + 3 = 6$ (no overflow)
 - $4 + 4 = 8$ (too big) $= 0$ (overflow)
 - $3 - 2 = 1$ (no overflow)
 - $3 - 4 = -1 = 7$ (underflow)



Integer Data Types

bits		minval	maxval
8	signed	-128	127
8	unsigned	0	255
16	signed	-32768	32767
16	unsigned	0	65535
32	signed	-2147483648	2147483647
32	unsigned	0	4294967295
64	signed	-9223372036854775808	9223372036854775807
64	unsigned	0	18446744073709551615



A number in a base b is represented as a sum of powers of b

- Base 10 (digits 0-9): $123 = 1 * 10^2 + 2 * 10^1 + 3 * 10^0$
- Base 2 (digits 0-1):
 $123 = 1 * 2^6 + 1 * 2^5 + 1 * 2^4 + 1 * 2^3 + 1 * 2^2 + 1 * 2^1 + 1 * 2^0$
- Base 8 (digits 0-7): $123 = 1 * 8^2 + 7 * 8^1 + 3 * 8^0$
- Base 16 (digits 0-9, A-F): $123 = 7 * 16^1 + 11 * 16^0$



Examples of Base numbers

- What is $(543)_8$ in decimal?
- What is $(101101)_2$ in decimal?
- What is $(3AB)_{16}$ in decimal?
- 78651 cannot be base 8, why not?



Base 16: Hexadecimal

Hex	bits	—	Hex	bits
0	0000		8	1000
1	0001		9	1001
2	0010		A	1010
3	0011		B	1011
4	0100		C	1100
5	0101		D	1101
6	0110		E	1110
7	0111		F	1111



Encoding a byte in Hexadecimal

- $D9 = 11011001$
- $AF = 10101111$
- $8C = 10001100$



Converting Between Binary and Decimal

2^7	2^6	2^5	2^4	2^3	2^2	2^1	2^0
128	64	32	16	8	4	2	1

- Method 1: sum powers of 2
- $10010010 = 128 + 16 + 2 = 146$
- Method 2: start from left
 - Start with 1
 - For each digit, multiply by 2
 - If the digit is 1, add 1
 - Example: $(((((1 * 2) * 2) * 2) + 1) * 2 * 2 * 2 + 1) * 2 = 146$



Integer Operations Overview

- Integer operations are fundamental in digital systems
- We'll cover arithmetic, logical, and shift operations
- Each operation will be demonstrated with 8-bit examples
- Understanding these operations is crucial for digital design



Addition

- Addition is performed bit by bit, carrying over when necessary
- Example (8-bit):

$$\begin{array}{r} 01101001 \\ + 00110110 \\ \hline 10011111 \end{array}$$

- Note: Overflow can occur if the result exceeds 8 bits



Subtraction

- Subtraction is often implemented as addition with two's complement
- Example (8-bit):

$$\begin{array}{r} 1\ 11 \\ 01101001 \\ 00110110 \\ \hline 00010011 \end{array}$$

- Two's complement of 00110110 is 11001010



Bitwise AND

- AND operation: 1 if both bits are 1, otherwise 0
- Example (8-bit):

$$\begin{array}{r} 01101001 \\ \& 00110110 \\ \hline 00110000 \end{array}$$

- Often used for masking specific bits



Bitwise OR

- OR operation: 1 if either bit is 1, otherwise 0
- Example (8-bit):

$$\begin{array}{r} 01101001 \\ | 00110110 \\ \hline 01111111 \end{array}$$

- Useful for setting specific bits



- XOR operation: 1 if bits are different, 0 if same
- Example (8-bit):

$$\begin{array}{r} 01101001 \\ \oplus 00110110 \\ \hline 01011111 \end{array}$$

- Often used in cryptography and error detection



Bitwise NOT

- Inverts all bits: 1 becomes 0, 0 becomes 1
- Example (8-bit):

$$\begin{array}{r} \sim 01101001 \\ \hline 10010110 \end{array}$$

- Useful for flipping specific bits



Logical Left Shift

- Shifts all bits to the left, filling rightmost with 0
- Example (8-bit, shift by 2):

$$\begin{array}{r} 01101001 \ll 2 \\ \hline 10100100 \end{array}$$

- Equivalent to multiplication by 2^n for n shifts



Logical Right Shift

- Shifts all bits to the right, filling leftmost with 0
- Example (8-bit, shift by 2):

$$\begin{array}{r} 01101001 \ll 2 \\ \hline 10100100 \end{array}$$

- Equivalent to multiplication by 2^n for n shifts (unsigned)



Rotate Left

- All bits move to the left
- Bits that drop off the end come in on right
- Example (8-bit, rotate by 2):

$$\begin{array}{r} 01101001 \text{ ROL } 2 \\ \hline 10100101 \end{array}$$

- Preserves all bits, useful in cryptography



Rotate Right

- Shifts all bits to the right, wrapping around
- Example (8-bit, rotate by 2):

$$\begin{array}{r} 01101001 \text{ ROR } 2 \\ \hline 10100101 \end{array}$$

- Example (8-bit, rotate by 2):
- Also preserves all bits, used in various algorithms



Arithmetic Left Shift

- Shifts bits to left just like logical shift
- If sign changes, it's an overflow
- Example (8-bit, shift by 2):

$$\begin{array}{r} 01101001 \lll 1 \\ \hline 11010010 \end{array}$$

- Equivalent to multiplication by 2^n for n shifts
- Can cause overflow if significant bits are shifted out



Arithmetic Right Shift

- Shift bits right preserving sign
- Leftmost bits are filled with the sign bit
- Example (8-bit, shift by 2):

$$\begin{array}{r} 11101001 \ggg 1 \\ \hline 11110100 \end{array}$$

- For positive numbers, equivalent to division by 2^n for n shifts
- For negative numbers, rounds towards negative infinity
- Preserves the sign of the number



Bitwise Operations in Verilog

- Verilog provides operators for various bitwise operations
- These operations are fundamental in digital design
- Examples (assuming 8-bit variables a and b):

```
c = a & b;           // Bitwise AND
d = a | b;           // Bitwise OR
e = a ^ b;           // Bitwise XOR
f = ~a;              // Bitwise NOT
g = a << 2;           // Logical left shift
h = a >> 2;           // Logical right shift
i = $signed(a) >>> 2; // Arithmetic right shift
j = {a[5:0], a[7:6]}; // Rotate left
k = {a[1:0], a[7:2]}; // Rotate right
```



- ASCII: American Standard Code for Information Interchange
- Maps characters to binary values
- ASCII table overview <https://www.ascii-code.com/>
- Example: 'A' = 01000001 = 65
- Example: 'B' = 01000010 = 66
- Example: 'a' = 01100001 = 97



Unicode UTF-8 Encoding

- Unicode is a standard for representing text of all languages
- Originally fit into 16 bits
- Now requires 19 bits because of emoji
- UTF-8 is a variable length encoding
- Useful when most of the text is English
- <https://symbl.cc/en/unicode-table/>
- <https://r12a.github.io/app-conversion/>

1st byte	2nd byte	3rd byte	4th byte	19-bit value
0xxxxxxx				0000000000000000xxxxxx
110yyyyy	10xxxxxx			0000000000yyyyyxxxxxx
1110zzzz	10yyzzzz	10yyyyyy	10xxxxxx	00000zzzzyyyyyyxxxxxx
11110uuu	10uuzzzz	10yyyyyy	10xxxxxx	uuuuuzzzzyyyyyyxxxxxx



Unicode UTF-8 Encoding

Let's decode the following bytes

00000000: 74 65 73 74 0a ce b5 cf 85 cf 87 ce b1 cf 81 ce b

00000010: cf 83 cf 84 cf 8e 0a E5 A4 A7

- 74 = 01110100 starts with 0, so look up ASCII
- same through 0A (ASCII for linefeed)
- CE B5 = 11001110 10110101 starts with 110, so it's a 2-byte char
- bits = 01110 110101 = 011 1011 0101 = 3B1
- Look up 3B1 in Unicode online (it's a Greek letter)
- E5 A4 A7 = 11100101 10100100 10100111 starts with 1110, so it's a 3-byte char
- bits = 0101 100100 100111 = 5 9 2 7
- Look up 0x5927 in Unicode online (Chinese)



Fixed Point Fractions

$$2^3 \mid 2^2 \mid 2^1 \mid 2^0 \mid . \mid 2^{-1} \mid 2^{-2}$$

- Fractions in binary are represented as negative powers
- $\frac{1}{2}, \frac{1}{4}, \frac{1}{8}$
- Examples

$$101.1 = 4 + 1 + \frac{1}{2} = 5.5$$

$$1.01 = 1 + \frac{1}{4} = 1.25$$

$$110.11 = 4 + 2 + \frac{1}{2} + \frac{1}{4}$$

$$1001.001 = 8 + 1 + \frac{1}{8}$$



Floating Point

- Fixed point represents fractions, but only a single size
- Floating point can represent values wildly different values
- IEEE-754

Single precision	32 bits
Double precision	64 bits
Quad Precision	128 bits (not yet in hardware)
Half Precision	16 bits (GPUs)
fp8	8 bits (GPUs)



R

◀ ◻ ▶ ◀ ◻ ▶ ◀ ≡ ▶ ◀ ≡ ▶ ≡ ↺ 🔍 ↻ 37/46

Unsigned Integers



Signed Integer Representations

-
-
-



American Standard Code for Information Interchange (ASCII)

- Represent characters with 7-bit codes
-
-
-
-



-
-
-
-
-



Fixed Point Representations

-
-
-
-
-



Floating Point Representation



Roundoff Errors



Limits of Floating Point



•
•

